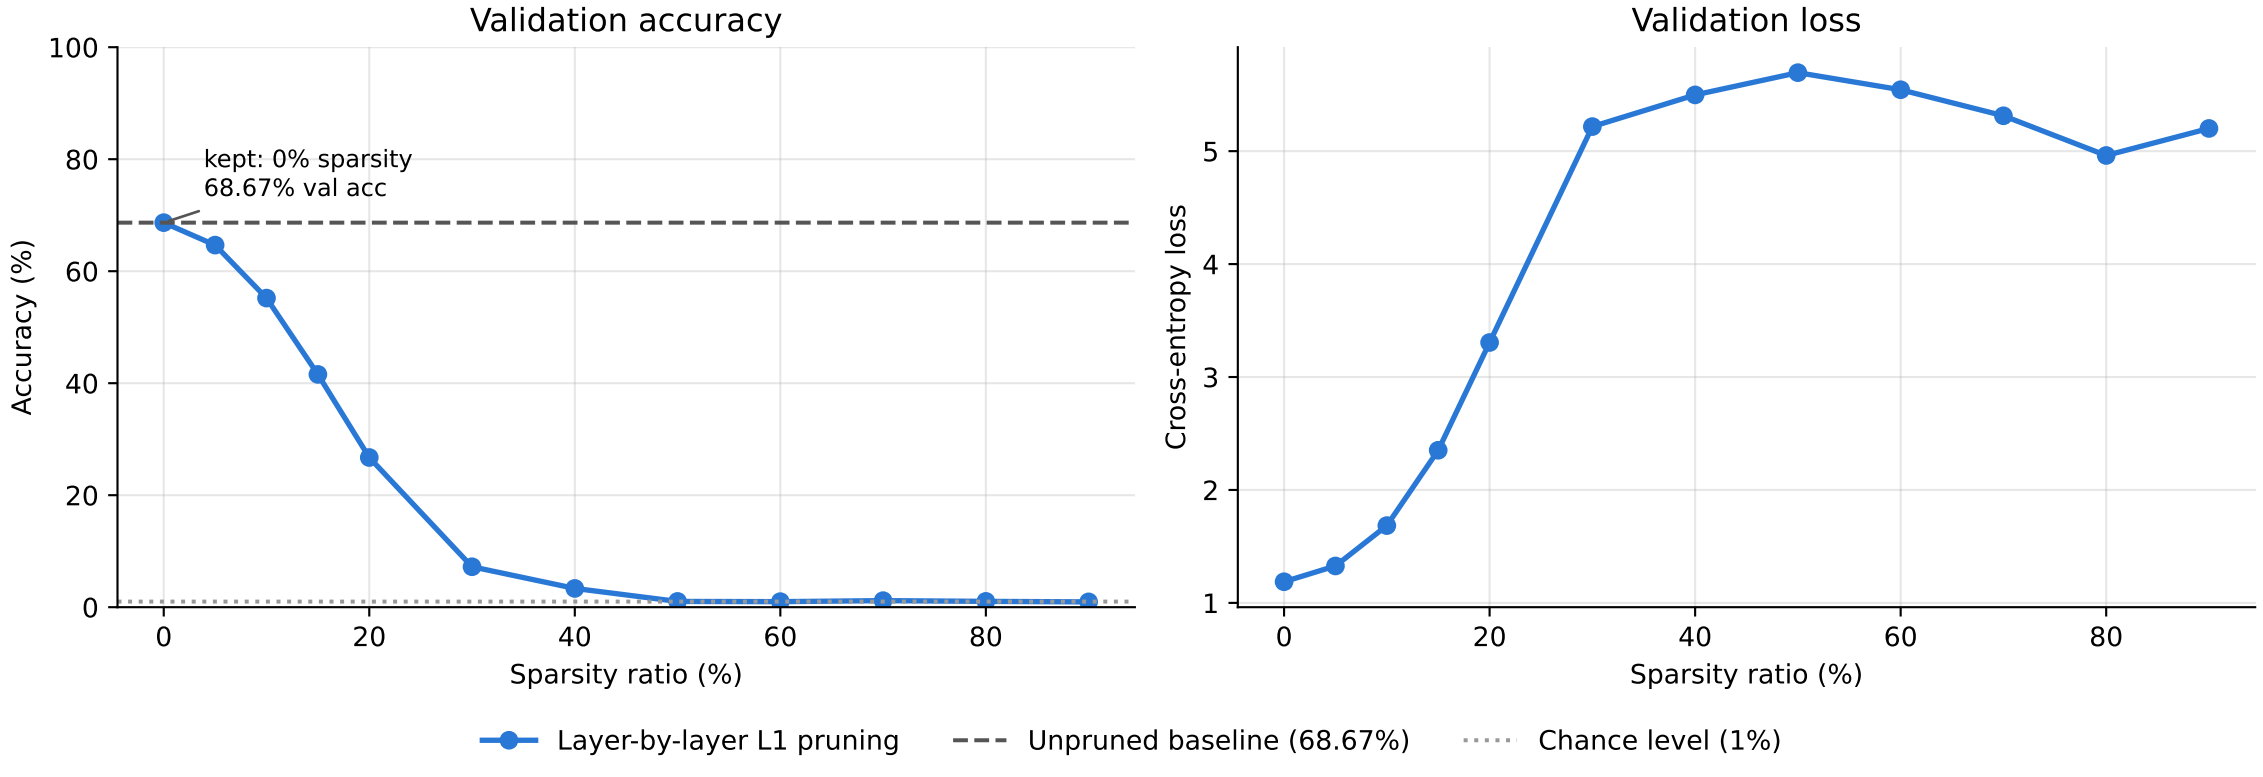
Exercise 2: accuracy / sparsity ratio curve (layer-by-layer pruning, no retraining)
Validation accuracy
Validation loss
kept: 0% sparsity
68.67% val acc
Accuracy (%)
Cross-entropy loss
Sparsity ratio (%)
Sparsity ratio (%)
Layer-by-layer L1 pruning
Unpruned baseline (68.67%)
Chance level (1%)